

```
In [1]: ▶ import pandas as pd
import numpy as np
from tensorflow import keras
from sklearn.model_selection import train_test_split
df= pd.read_csv('df-filteredfinalN.csv')
from keras.losses import categorical_crossentropy
from keras.utils import np_utils
from keras.models import Sequential
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.layers import Dense
from keras.layers.core import Dropout, Activation, Flatten
import seaborn as sns
import matplotlib.pyplot as plt
import timeit
```

In this part of the solution, different Convolutional Neural Network(CNN) models are created for the training data from the Fashion MNIST dataset[1]. The model is then tested on the test set and the accuracy is then measured. All the important block codes for creating the models are titled and explained below.

Forming the Dataframe 'df'

The block below was used to load the training x(features) dataset and training y(label) dataset named 'trainX.csv' and 'trainy.csv' respectively. Then, the 'Id' column present in the training x dataset is removed. Both datasets are concatenated to form the 'df' dataframe.

In [2]:

```
cols = list(pd.read_csv("trainX.csv", nrows =1))

xt=pd.read_csv('trainX.csv', usecols = [i for i in cols if i != 'Id'])

yt=pd.read_csv('trainy.csv')

df=pd.concat([xt,yt], axis=1)

del df['Id']

xtest=pd.read_csv('testX.csv')

ytest=pd.read_csv('testy.csv')

del xtest['Id']
del ytest['Id']

df.describe()
```


Out[2]:

	0	1	2	3	4	5
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
mean	0.000483	0.006067	0.033700	0.090450	0.241383	0.392517
std	0.061507	0.292286	1.178115	2.306445	4.365543	5.702057
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	14.000000	45.000000	218.000000	185.000000	227.000000	229.000000

8 rows × 785 columns


Dividing Features and Encoding Labels

The block below is used to separate the dependent and independent features. The labels are also encoded using one hot coding.

```
In [3]:  #Separating dependent and independent features  
  
x =df.loc[:, '0': '783'].values.reshape(len(df),28,28,1)  
  
y1 = df[' Label'].values  
  
#One hot ecoding for label  
  
y = keras.utils.to_categorical(y1, num_classes=5)
```

Splitting Train and Test Sets

The Train and Test sets are separated with the test size being 0.2.

```
In [4]:  x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2,randc  
x_train=x_train.astype('float32')  
xtest=xtest.astype('float32')
```

Model 1: 2 Convolutional Layers with 32 Filters(ReLU Activation)

The model created in the block below contains 2 convolutional layers with 32 filter, ReLU activation and 'Same' padding.

```

In [44]: model1 = Sequential()

model1.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',
                  activation = 'relu', input_shape = (28,28,1)))
model1.add(MaxPooling2D(pool_size=(2,2)))

model1.add(Conv2D(filters = 32, kernel_size = (3,3),padding = 'Same',
                  activation = 'relu'))
model1.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))

model1.add(Flatten())
model1.add(Dense(150, activation = "relu"))
model1.add(Dense(5, activation = "softmax"))

model1.compile(optimizer = keras.optimizers.Adam() , loss = "categorical_crossentropy")

start=timeit.default_timer()
classifier1=model1.fit(x_train,y_train,batch_size = 32,epochs = 15,verbose=1,
stop=timeit.default_timer())

```

```

Epoch 1/15
1005/1005 [=====] - 67s 64ms/step - loss: 0.9458 -
accuracy: 0.6745 - val_loss: 0.6141 - val_accuracy: 0.7429
Epoch 2/15
1005/1005 [=====] - 64s 64ms/step - loss: 0.5618 -
accuracy: 0.7579 - val_loss: 0.5494 - val_accuracy: 0.7658
Epoch 3/15
1005/1005 [=====] - 69s 68ms/step - loss: 0.5015 -
accuracy: 0.7865 - val_loss: 0.5002 - val_accuracy: 0.7896
Epoch 4/15
1005/1005 [=====] - 64s 63ms/step - loss: 0.4502 -
accuracy: 0.8078 - val_loss: 0.5001 - val_accuracy: 0.7902
Epoch 5/15
1005/1005 [=====] - 68s 68ms/step - loss: 0.4129 -
accuracy: 0.8230 - val_loss: 0.4725 - val_accuracy: 0.8032
Epoch 6/15
1005/1005 [=====] - 70s 69ms/step - loss: 0.3861 -
accuracy: 0.8363 - val_loss: 0.5028 - val_accuracy: 0.8093
Epoch 7/15
1005/1005 [=====] - 72s 72ms/step - loss: 0.3611 -
accuracy: 0.8462 - val_loss: 0.4950 - val_accuracy: 0.8044
Epoch 8/15
1005/1005 [=====] - 52s 52ms/step - loss: 0.3362 -
accuracy: 0.8605 - val_loss: 0.5078 - val_accuracy: 0.8094
Epoch 9/15
1005/1005 [=====] - 36s 35ms/step - loss: 0.3130 -
accuracy: 0.8696 - val_loss: 0.4860 - val_accuracy: 0.8205
Epoch 10/15
1005/1005 [=====] - 37s 36ms/step - loss: 0.2906 -

```

```
accuracy: 0.8801 - val_loss: 0.5109 - val_accuracy: 0.8180
Epoch 11/15
1005/1005 [=====] - 37s 37ms/step - loss: 0.2651 -
accuracy: 0.8909 - val_loss: 0.6107 - val_accuracy: 0.8008
Epoch 12/15
1005/1005 [=====] - 37s 37ms/step - loss: 0.2546 -
accuracy: 0.8954 - val_loss: 0.5514 - val_accuracy: 0.8143
Epoch 13/15
1005/1005 [=====] - 34s 34ms/step - loss: 0.2341 -
accuracy: 0.9049 - val_loss: 0.5930 - val_accuracy: 0.8147
Epoch 14/15
1005/1005 [=====] - 37s 37ms/step - loss: 0.2282 -
accuracy: 0.9080 - val_loss: 0.6199 - val_accuracy: 0.8061
Epoch 15/15
1005/1005 [=====] - 38s 38ms/step - loss: 0.2140 -
accuracy: 0.9140 - val_loss: 0.6658 - val_accuracy: 0.8061
```

Model 1 Training Accuracy:

The block below trains Model 1 and prints the total training time, the training accuracy and the validation accuracy.

```
In [55]: ▶ print('Max. Accuracy for the Training Set:')
          print(max(classifier1.history['accuracy']))

          print('Max. Accuracy for the Validation Set:')
          print(max(classifier1.history['val_accuracy']))
```

```
Max. Accuracy for the Training Set:
0.9146117568016052
Max. Accuracy for the Validation Set:
0.8152894377708435
```

A 91.46 percent of accuracy was achieved in the training set and a 81.53 percent of accuracy was achieved in the validation set.

Training Loss vs. Training Epoch Plot

```
In [19]: y= classifier1.history['loss']
x= range(0,15)
y1=classifier1.history['val_loss']

plt.plot(x,y,label='Training Loss')
plt.plot(x,y1,'g',label='Validation Loss')

plt.ylabel('Training Loss')
plt.xlabel('Training Epoch')
plt.title('Training Loss Vs. Training Epoch Plot ')
plt.legend(['Training','Validation'])
```

Out[19]: <matplotlib.legend.Legend at 0x23386af4fa0>



Figure 1. Training Loss Vs. Training Epoch Plot

Classification Accuracy Vs. Training Epoch Plot

```
In [20]: y= classifier1.history['accuracy']
y1= classifier1.history['val_accuracy']
x= range(0,15)

plt.plot(x,y, label='Training Accuracy')
plt.plot(x,y1, 'g', label='Validation Accuracy')

plt.ylabel('Classification Accuracy')
plt.xlabel('Training Epoch')
plt.title('Classification Accuracy Vs. Training Epoch Plot ')
plt.legend(['Training','Validation'])
```

Out[20]: <matplotlib.legend.Legend at 0x23386b36b80>




Figure 2. Classification Accuracy Vs. Training Epoch Plot

Model 2: 2 Convolutional Layers with 16 Filters(ReLU Activation), Valid Padding, and Dropout(0.5)

The model created in the block below contains 2 convolutional layers with 16 filter, ReLU activation and 'Valid' padding[2] with a dropout of 0.5.

Model 2 Training Accuracy:

```
In [48]: model2 = Sequential()

model2.add(Conv2D(filters = 16, kernel_size = (5,5),padding = 'Valid',
                  activation = 'relu', input_shape = (28,28,1)))
model2.add(MaxPooling2D(pool_size=(2,2)))

model2.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Valid',
                  activation = 'relu'))
model2.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))

model2.add(Flatten())
model2.add(Dense(150, activation = "relu"))
model2.add(Dense(5, activation = "softmax"))

model2.compile(optimizer = keras.optimizers.Adam() , loss = "categorical_crossentropy")

classifier2=model2.fit(x_train,y_train,batch_size = 32,epochs = 20,verbose=1,
```

```
Epoch 1/20
1005/1005 [=====] - 29s 26ms/step - loss: 1.2856 - accuracy: 0.6179 - val_loss: 0.7287 - val_accuracy: 0.6839
Epoch 2/20
1005/1005 [=====] - 33s 32ms/step - loss: 0.6838 - accuracy: 0.7057 - val_loss: 0.6547 - val_accuracy: 0.7213
Epoch 3/20
1005/1005 [=====] - 38s 38ms/step - loss: 0.5973 - accuracy: 0.7450 - val_loss: 0.5962 - val_accuracy: 0.7493
Epoch 4/20
1005/1005 [=====] - 36s 36ms/step - loss: 0.5474 - accuracy: 0.7670 - val_loss: 0.5743 - val_accuracy: 0.7569
Epoch 5/20
1005/1005 [=====] - 33s 33ms/step - loss: 0.5120 - accuracy: 0.7815 - val_loss: 0.6163 - val_accuracy: 0.7366
Epoch 6/20
1005/1005 [=====] - 36s 35ms/step - loss: 0.4799 - accuracy: 0.7978 - val_loss: 0.5422 - val_accuracy: 0.7763
Epoch 7/20
1005/1005 [=====] - 36s 36ms/step - loss: 0.4578 - accuracy: 0.8061 - val_loss: 0.5591 - val_accuracy: 0.7700
Epoch 8/20
1005/1005 [=====] - 42s 42ms/step - loss: 0.4373 - accuracy: 0.8153 - val_loss: 0.5829 - val_accuracy: 0.7606
Epoch 9/20
1005/1005 [=====] - 43s 43ms/step - loss: 0.4181 - accuracy: 0.8244 - val_loss: 0.5922 - val_accuracy: 0.7597
Epoch 10/20
1005/1005 [=====] - 41s 41ms/step - loss: 0.3961 - accuracy: 0.8328 - val_loss: 0.5764 - val_accuracy: 0.7787
Epoch 11/20
```



```

1005/1005 [=====] - 40s 40ms/step - loss: 0.3799 -
accuracy: 0.8419 - val_loss: 0.6003 - val_accuracy: 0.7722
Epoch 12/20
1005/1005 [=====] - 41s 40ms/step - loss: 0.3658 -
accuracy: 0.8464 - val_loss: 0.6132 - val_accuracy: 0.7845
Epoch 13/20
1005/1005 [=====] - 40s 40ms/step - loss: 0.3497 -
accuracy: 0.8521 - val_loss: 0.6889 - val_accuracy: 0.7708
Epoch 14/20
1005/1005 [=====] - 40s 40ms/step - loss: 0.3471 -
accuracy: 0.8554 - val_loss: 0.6349 - val_accuracy: 0.7801
Epoch 15/20
1005/1005 [=====] - 41s 40ms/step - loss: 0.3271 -
accuracy: 0.8613 - val_loss: 0.6290 - val_accuracy: 0.7880
Epoch 16/20
1005/1005 [=====] - 40s 40ms/step - loss: 0.3255 -
accuracy: 0.8655 - val_loss: 0.7700 - val_accuracy: 0.7529
Epoch 17/20
1005/1005 [=====] - 40s 40ms/step - loss: 0.3139 -
accuracy: 0.8684 - val_loss: 0.6969 - val_accuracy: 0.7705
Epoch 18/20
1005/1005 [=====] - 40s 40ms/step - loss: 0.3007 -
accuracy: 0.8759 - val_loss: 0.7729 - val_accuracy: 0.7766
Epoch 19/20
1005/1005 [=====] - 41s 41ms/step - loss: 0.2885 -
accuracy: 0.8795 - val_loss: 0.7508 - val_accuracy: 0.7875
Epoch 20/20
1005/1005 [=====] - 42s 42ms/step - loss: 0.2776 -
accuracy: 0.8852 - val_loss: 0.8767 - val_accuracy: 0.7679

```

The block below prints the total training time, the training accuracy and the validation accuracy.

In [8]: ▶

```

print('Max. Accuracy for the Training Set:')
print(max(classifier2.history['accuracy']))

print('Max. Accuracy for the Validation Set:')
print(max(classifier2.history['val_accuracy']))

```

```

Max. Accuracy for the Training Set:
0.8856307864189148
Max. Accuracy for the Validation Set:
0.7944574356079102

```

Training Loss vs. Training Epoch Plot

```
In [51]: y= classifier2.history['loss']
x= range(0,20)
y1=classifier2.history['val_loss']

plt.plot(x,y,label='Training Loss')
plt.plot(x,y1,'g',label='Validation Loss')

plt.ylabel('Training Loss')
plt.xlabel('Training Epoch')
plt.title('Training Loss Vs. Training Epoch Plot ')
plt.legend(['Training','Validation'])
```

Out[51]: <matplotlib.legend.Legend at 0x22589d73700>




Figure 3. Training Loss Vs. Training Epoch Plot

Classification Accuracy Vs. Training Epoch Plot

```
In [52]: y= classifier2.history['accuracy']
y1= classifier2.history['val_accuracy']
x= range(0,20)

plt.plot(x,y, label='Training Accuracy')
plt.plot(x,y1, 'g', label='Validation Accuracy')

plt.ylabel('Classification Accuracy')
plt.xlabel('Training Epoch')
plt.title('Classification Accuracy Vs. Training Epoch Plot ')
plt.legend(['Training', 'Validation'])
```

Out[52]: <matplotlib.legend.Legend at 0x2258965edf0>




Figure 4. Classification Accuracy Vs. Training Epoch Plot

Model 3: 4 Convolutional Layers with 16 filters(layer 1,2), 32 Filters(Layer 3,4) with (ReLU Activation), and Valid Padding

```

In [29]: model3 = Sequential()

model3.add(Conv2D(filters = 16, kernel_size = (5,5),padding = 'Same',
                  activation = 'relu', input_shape = (28,28,1)))
model3.add(MaxPooling2D(pool_size=(2,2)))

model3.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',
                  activation = 'relu'))
model3.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))

model3.add(Conv2D(filters = 32, kernel_size = (3,3),padding = 'Same',
                  activation = 'relu'))
model3.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))

model3.add(Conv2D(filters = 32, kernel_size = (3,3),padding = 'Same',
                  activation = 'relu'))
model3.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))

model3.add(Flatten())
model3.add(Dense(150, activation = "relu"))
model3.add(Dense(5, activation = "softmax"))

model3.compile(optimizer = keras.optimizers.Adam() , loss = "categorical_crossentropy")

start=timeit.default_timer()
classifier3=model3.fit(x_train,y_train,batch_size = 32,epochs = 20,verbose=1,
stop=timeit.default_timer())

```

```

Epoch 1/20
1005/1005 [=====] - 23s 23ms/step - loss: 0.803
6 - accuracy: 0.6566 - val_loss: 0.6461 - val_accuracy: 0.7227
Epoch 2/20
1005/1005 [=====] - 30s 30ms/step - loss: 0.613
6 - accuracy: 0.7373 - val_loss: 0.6182 - val_accuracy: 0.7426
Epoch 3/20
1005/1005 [=====] - 43s 42ms/step - loss: 0.540
9 - accuracy: 0.7681 - val_loss: 0.5585 - val_accuracy: 0.7653
Epoch 4/20
1005/1005 [=====] - 57s 56ms/step - loss: 0.501
2 - accuracy: 0.7883 - val_loss: 0.5005 - val_accuracy: 0.7864
Epoch 5/20
1005/1005 [=====] - 51s 50ms/step - loss: 0.467
3 - accuracy: 0.8015 - val_loss: 0.5767 - val_accuracy: 0.7482
Epoch 6/20
1005/1005 [=====] - 50s 50ms/step - loss: 0.450
4 - accuracy: 0.8099 - val_loss: 0.4658 - val_accuracy: 0.8043
Epoch 7/20
1005/1005 [=====] - 62s 62ms/step - loss: 0.425
6 - accuracy: 0.8195 - val_loss: 0.4749 - val_accuracy: 0.8068
Epoch 8/20

```

```

1005/1005 [=====] - 63s 62ms/step - loss: 0.403
5 - accuracy: 0.8296 - val_loss: 0.4608 - val_accuracy: 0.8094
Epoch 9/20
1005/1005 [=====] - 60s 59ms/step - loss: 0.393
7 - accuracy: 0.8362 - val_loss: 0.4950 - val_accuracy: 0.7985
Epoch 10/20
1005/1005 [=====] - 59s 59ms/step - loss: 0.380
2 - accuracy: 0.8415 - val_loss: 0.4634 - val_accuracy: 0.8094
Epoch 11/20
1005/1005 [=====] - 60s 60ms/step - loss: 0.372
6 - accuracy: 0.8463 - val_loss: 0.4617 - val_accuracy: 0.8111
Epoch 12/20
1005/1005 [=====] - 59s 59ms/step - loss: 0.361
4 - accuracy: 0.8483 - val_loss: 0.4540 - val_accuracy: 0.8150
Epoch 13/20
1005/1005 [=====] - 60s 60ms/step - loss: 0.349
9 - accuracy: 0.8536 - val_loss: 0.4668 - val_accuracy: 0.8193
Epoch 14/20
1005/1005 [=====] - 59s 58ms/step - loss: 0.345
1 - accuracy: 0.8567 - val_loss: 0.4635 - val_accuracy: 0.8120
Epoch 15/20
1005/1005 [=====] - 63s 63ms/step - loss: 0.325
5 - accuracy: 0.8641 - val_loss: 0.4467 - val_accuracy: 0.8208
Epoch 16/20
1005/1005 [=====] - 41s 40ms/step - loss: 0.329
3 - accuracy: 0.8613 - val_loss: 0.4577 - val_accuracy: 0.8190s: 0.3293
- accuracy: 0.86
Epoch 17/20
1005/1005 [=====] - 36s 36ms/step - loss: 0.316
3 - accuracy: 0.8690 - val_loss: 0.4678 - val_accuracy: 0.8195
Epoch 18/20
1005/1005 [=====] - 36s 36ms/step - loss: 0.310
5 - accuracy: 0.8707 - val_loss: 0.4687 - val_accuracy: 0.8159
Epoch 19/20
1005/1005 [=====] - 36s 36ms/step - loss: 0.312
2 - accuracy: 0.8702 - val_loss: 0.4735 - val_accuracy: 0.8171
Epoch 20/20
1005/1005 [=====] - 36s 36ms/step - loss: 0.300
9 - accuracy: 0.8742 - val_loss: 0.4989 - val_accuracy: 0.8070

```

Model 3 Training Accuracy:

The block below prints the total training time, the training accuracy and the validation accuracy.

In [32]:

```
print('Max. Accuracy for the Training Set:')
print(max(classifier3.history['accuracy']))

print('Max. Accuracy for the Validation Set:')
print(max(classifier3.history['val_accuracy']))
```

```
Max. Accuracy for the Training Set:
0.8741876482963562
Max. Accuracy for the Validation Set:
0.8207815289497375
```

Training Loss vs. Training Epoch Plot

In [33]:

```
y= classifier3.history['loss']
x= range(0,20)
y1=classifier3.history['val_loss']

plt.plot(x,y,label='Training Loss')
plt.plot(x,y1,'g',label='Validation Loss')

plt.ylabel('Training Loss')
plt.xlabel('Training Epoch')
plt.title('Training Loss Vs. Training Epoch Plot ')
plt.legend(['Training','Validation'])
```

Out[33]: <matplotlib.legend.Legend at 0x23386c2d8e0>

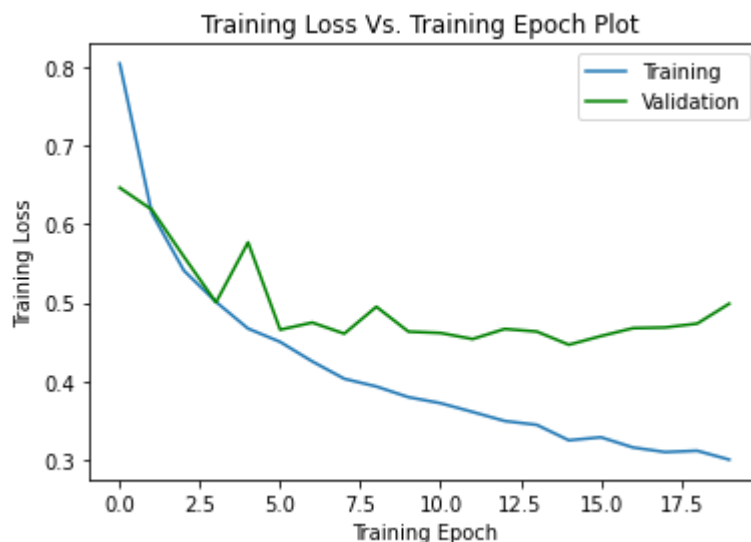


Figure 5. Training Loss Vs. Training Epoch Plot

Classification Accuracy Vs. Training Epoch

Plot

```
In [34]: ▶ y= classifier3.history['accuracy']
y1= classifier3.history['val_accuracy']
x= range(0,20)

plt.plot(x,y, label='Training Accuracy')
plt.plot(x,y1,'g', label='Validation Accuracy')

plt.ylabel('Classification Accuracy')
plt.xlabel('Training Epoch')
plt.title('Classification Accuracy Vs. Training Epoch Plot ')
plt.legend(['Training', 'Validation'])
```

Out[34]: <matplotlib.legend.Legend at 0x23386e148e0>




Figure 6. Classification Accuracy Vs. Training Epoch Plot

Review

1. Runtime Performance:

A. Model 1: The maximum accuracy for the training set was 91.46% and the maximum accuracy for the validation set was 81.53%. On the 'Training loss Vs. Training Epoch' plot, the training loss error reached 0.2(20%) after 12 epochs. For the validation, the training loss error reached 0.65(65%) after 14 epochs as shown in figure 1. On the 'Classification Accuracy Vs. Training Epoch Plot', the validation set curve settled at 0.78 epoch(78%) and the training curve reached almost 0.94(94%) in classification accuracy after about 17 epochs as shown in the figure 2.

B. Model 2: The maximum accuracy for the training set was 88.52% and the maximum accuracy for the validation set was 78.8%. On the 'Training loss Vs. Training Epoch' plot, the training loss error settled and remained below 0.4(40%) after 17 epochs. For the validation set, the training loss error went below 0.8 after 1 epoch but then came back up after 17 epochs as shown in figure 3. On the 'Classification Accuracy Vs. Training Epoch Plot', the validation set curve settled at 0.75 epoch(75%) and the training curve reached almost 0.9(90%) in classification accuracy after about 17 epochs as shown in the figure 4.

C. Model 3: The maximum accuracy for the training set was 87.42% and the maximum accuracy for the validation set was 82.08%. On the 'Training loss Vs. Training Epoch' plot, the training loss error reached 0.3(30%) after 17 epochs. For the validation set, the training loss error reached and remained 0.47(47%) shown in figure 5. On the 'Classification Accuracy Vs. Training Epoch Plot', the validation set curve remained at 0.8(80%) and the training curve reached 0.9(90%) in classification accuracy after about 2 epochs as shown in the figure 6.

2. Comparison: Model 1 had the best performance, followed by Model 3. Model 2 had the worst performance as it only had 78.8% maximum accuracy in the validation dataset.

References

[1]<https://www.datacamp.com/community/tutorials/convolutional-neural-networks-python#predictions> (<https://www.datacamp.com/community/tutorials/convolutional-neural-networks-python#predictions>)

[2]<https://www.geeksforgeeks.org/types-of-padding-in-convolution-layer/> (<https://www.geeksforgeeks.org/types-of-padding-in-convolution-layer/>)